

IoT Sensor Data Analysis Challenge

Problem Statement

Objective: As an AI intern working on an IoT project, your task is to analyze temperature data from a sensor in a controlled environment to identify recurring temporal patterns using machine learning. The dataset is provided as a CSV file with columns `device_id`, `ts` (Unix timestamp in milliseconds), and `temperature`. You will model the expected temperature behavior, detect deviations from these patterns, and implement alerting rules for significant anomalies.

Tasks

1. Data Loading and Preprocessing

- Load the CSV file into a data structure.
- Convert the `ts` column from Unix timestamps (milliseconds) to a human-readable format: `YYYY-MM-DD HH:MM:SS`.
- Handle missing or invalid data appropriately (e.g., imputation, removal, or flagging).

2. Pattern Analysis with Machine Learning

- Use a machine learning model to capture periodic temperature patterns (e.g., cycles over minutes or hours).
- Extract insights from the model, such as:
 - Cycle period (e.g., time between peaks).
 - Amplitude (e.g., range of temperature fluctuations).
- Report model performance using metrics like Mean Squared Error (MSE) and R^2 score.

3. Deviation Detection

- Identify significant deviations by comparing actual and predicted temperatures (e.g., `|actual - predicted| > 2°C`).
- Add an `is_deviation` column to the dataset:
 - `1` for records with deviations.
 - `0` otherwise.

4. Visualization

- Create a time-series plot showing:
 - Actual temperatures.
 - Predicted temperatures.
 - Markers for deviations.

- Plot the cyclic pattern (e.g., predicted temperature over one or more cycles).
- Generate a histogram of prediction errors, highlighting the deviation threshold (e.g., 2°C).
- Save all plots as PNG files:
 - `timeseries.png`
 - `cycle_pattern.png`
 - `error_histogram.png`

5. Alerting Rules

- Define rules for generating alerts based on deviations (e.g., threshold-based or consecutive deviations).
- Summarize alerts in a report or table, including:
 - Timestamp.
 - Actual temperature.
 - Predicted temperature.
 - Reason for the alert (e.g., deviation magnitude).

6. Output and Report

- Save a CSV file for each temperature dataset with the following columns:
 - `device_id`
 - Reformatted `ts` (YYYY-MM-DD HH:MM:SS)
 - `temperature`
 - `predicted_temperature`
 - `is_deviation`
- Provide a summary report (e.g., in a markdown cell, document, or code comments) that includes:
 - Total number of records processed.
 - Key patterns identified (e.g., cycle period, temperature range).
 - Number of deviations detected and alerts generated.
 - Model performance metrics (e.g., MSE, R^2).
 - Comparison of patterns between the two temperature datasets (e.g., differences in cycle period or amplitude).

7. Bonus: Motor Current Analysis (Optional)

- Apply your machine learning method to the motor current dataset (CSV with columns `device_id`, `ts`, `motor_current` in amperes).
- In a markdown cell or document, discuss:
 - How well the method performed on the current data compared to the temperature data.
 - Any challenges encountered (e.g., noisier data, different periodicity).
 - Differences in patterns between temperature and current data.

Requirements

- **Language:** Python is preferred, but other languages are acceptable.
 - **Python:** Submit as a Jupyter notebook.
 - **Other languages:** Submit a well-documented script with a separate report document.
- **Input:**
 - Temperature CSV 1 (temperature-data-device-1.csv)
 - Temperature CSV 2 (temperature-data-device-2.csv)
 - Motor current CSV (motor-data-device-3.csv)
 - All CSVs contain: **device_id**, **ts** (Unix timestamp in milliseconds), and either **temperature** (°C) or **motor_current** (amperes).
- **Output:**
 - CSV files for each dataset with the specified columns.
 - Image files for visualizations.
 - A Jupyter notebook (Python) or script (other languages) with clear documentation and implementation.
- The use of AI tools is permitted, but you must disclose in your report how they were used and for what purposes (e.g., code generation, analysis).